

I would use the following **30 references**.

References

- [1] A. Janosi, W. Steinbrunn, M. Pfisterer, and R. Detrano, "Heart Disease," UCI Machine Learning Repository, 1989. [Online]. Available: [UCI Heart Disease Dataset](#)
 - [2] B. Strack *et al.*, "Impact of HbA1c Measurement on Hospital Readmission Rates: Analysis of 70,000 Clinical Database Patient Records," *BioMed Research International*, vol. 2014, Art. no. 781670, 2014.
 - [3] L. Rubini, P. Soundarapandian, and P. Eswaran, "Chronic Kidney Disease," UCI Machine Learning Repository, 2015. [Online]. Available: [UCI Chronic Kidney Disease Dataset](#)
 - [4] B. Ramana and N. Venkateswarlu, "ILPD (Indian Liver Patient Dataset)," UCI Machine Learning Repository, 2012. [Online]. Available: [UCI ILPD Dataset](#)
 - [5] W. H. Wolberg, O. L. Mangasarian, N. Street, and W. H. Street, "Breast Cancer Wisconsin (Diagnostic)," UCI Machine Learning Repository, 1995. [Online]. Available: [UCI Breast Cancer Wisconsin Dataset](#)
 - [6] H. A. Little *et al.*, "Parkinsons," UCI Machine Learning Repository, 2007. [Online]. Available: [UCI Parkinsons Dataset](#)
 - [7] R. S. Gong, "Hepatitis," UCI Machine Learning Repository, 1988. [Online]. Available: [UCI Hepatitis Dataset](#)
 - [8] D. Chicco and G. Jurman, "Machine learning can predict survival of patients with heart failure from serum creatinine and ejection fraction alone," *BMC Medical Informatics and Decision Making*, vol. 20, Art. no. 16, 2020, doi: 10.1186/s12911-020-1023-5. ([PubMed Central \(PMC\)](#))
 - [9] fedesoriano, "Stroke Prediction Dataset," Kaggle, 2021. [Online]. Available: [Kaggle Stroke Prediction Dataset](#)
-

Machine Learning

- [10] L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001, doi: 10.1023/A:1010933404324.
- [11] C. Cortes and V. Vapnik, "Support-vector networks," *Machine Learning*, vol. 20, pp. 273–297, 1995, doi: 10.1007/BF00994018.

[12] T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proc. 22nd ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining*, 2016, pp. 785–794, doi: 10.1145/2939672.2939785.

[13] D. W. Hosmer, S. Lemeshow, and R. X. Sturdivant, *Applied Logistic Regression*, 3rd ed. Hoboken, NJ, USA: Wiley, 2013.

[14] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, “SMOTE: Synthetic minority over-sampling technique,” *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002.

[15] D. Chicco, N. Tötsch, and G. Jurman, “The Matthews correlation coefficient (MCC) is more reliable than F1 score and accuracy in binary classification evaluation,” *BioData Mining*, vol. 14, Art. no. 13, 2021.

Explainable AI

[16] S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” in *Advances in Neural Information Processing Systems*, vol. 30, 2017.

[17] J. Jung, H. Lee, H. Jung, and H. Kim, “Essential properties and explanation effectiveness of explainable artificial intelligence in healthcare: A systematic review,” *Heliyon*, vol. 9, no. 5, 2023, Art. no. e16110.

[18] N. A. Aziz *et al.*, “Explainable AI in healthcare: Systematic review of clinical decision support systems,” *medRxiv*, 2024. [Online]. Available: [medRxiv record](#)

Note: [18] is a preprint. Do **not** write it as a peer-reviewed journal paper.

RAG and Medical LLMs

[19] P. Lewis *et al.*, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Advances in Neural Information Processing Systems*, vol. 33, 2020. ([NeurIPS Proceedings](#))

[20] A. J. Thirunavukarasu *et al.*, “Large language models in medicine,” *Nature Medicine*, vol. 29, pp. 1930–1940, 2023, doi: 10.1038/s41591-023-02448-8.

[21] H. Zhou *et al.*, “A survey of large language models in medicine: Progress, application, and challenge,” 2023, arXiv:2311.05112.

[22] H. Nori *et al.*, “Capabilities of GPT-4 on medical challenge problems,” Microsoft Research, 2023. [Online]. Available: [Microsoft Research paper record](#)

[23] K. Singhal *et al.*, “Towards expert-level medical question answering with large language models,” 2023, arXiv:2305.09617.

[24] L. Wang *et al.*, “Applications and concerns of ChatGPT and other conversational large language models in health care: Systematic review,” *Journal of Medical Internet Research*, vol. 26, 2024.

Multi-Agent Medical AI

[25] X. Tang, A. Zou, Z. Zhang, Z. Li, Y. Zhao, X. Zhang, A. Cohan, and M. Gerstein, “MedAgents: Large language models as collaborators for zero-shot medical reasoning,” in *Findings of the Association for Computational Linguistics: ACL 2024*, Bangkok, Thailand, 2024, pp. 599–621, doi: 10.18653/v1/2024.findings-acl.33. ([ACL Anthology](#))

[26] Y. Kim, C. Park, H. Jeong, Y. S. Chan, X. Xu, D. McDuff, H. Lee, M. Ghassemi, C. Breazeal, and H. W. Park, “MDAgents: An adaptive collaboration of LLMs for medical decision-making,” in *Advances in Neural Information Processing Systems* 37, 2024, pp. 79410–79452, doi: 10.52202/079017-2522. ([Proceedings](#))

[27] S. Delourme *et al.*, “Measured performance and healthcare professional perception of large language models used as clinical decision support systems: A scoping review,” *Studies in Health Technology and Informatics*, 2024, doi: 10.3233/SHTI240543.

Healthcare Interoperability / FHIR

[28] K. Bender and K. Sartipi, “HL7 FHIR: An agile and RESTful approach to healthcare information exchange,” in *Proc. IEEE International Conference on Healthcare Informatics*, 2013.

[29] A. Riquelme, P. Costa, and C. Martinez, “Large language models for automating clinical data standardization: HL7 FHIR use case,” 2025, arXiv:2507.03067.

Recent Biomedical AI

[30] H. Yu *et al.*, “Large language models in biomedical and health informatics: A review with bibliometric analysis,” 2024, arXiv:2403.16303.

Very important for your paper

For your **MediGenie** project, don't simply paste all 30 references into the end and leave them unused.

Use them like this:

- [1]–[9] → your nine datasets
- [10]–[15] → your ML algorithms and evaluation
- [16]–[18] → SHAP/XAI
- [19] → RAG
- [20]–[24] → medical LLM background
- [25]–[27] → your multi-agent architecture
- [28]–[29] → FHIR/API interoperability
- [30] → broader biomedical AI background

And [8], [25], and [26] are particularly important for your comparison/related-work section: [8] is a direct Heart Failure ML comparison, while [25] and [26] support the multi-agent medical reasoning component. ([PubMed Central \(PMC\)](#))

One correction I strongly recommend

Your paper should describe [25] **MedAgents as the 2024 ACL Findings publication**, because that is the finalized peer-reviewed conference version. The earlier 2023 arXiv version exists, but the ACL 2024 version is the better reference to put in an IEEE bibliography. ([ACL Anthology](#))

Similarly, [19] is **officially a NeurIPS 2020 paper**, not merely an arXiv paper. ([NeurIPS Proceedings](#))

Yes. I would use a **30-reference bibliography**, but I would not make all 30 “recent” papers. Your paper needs a mixture of **dataset provenance + foundational ML methods + recent medical AI/multi-agent/RAG/FHIR research**.

I verified the recent references against publisher/official research pages and DOI records. For example, MDAgents is a NeurIPS 2024 paper with DOI 10.52202/079017-2522; the original MedAgents work is documented as a 2023 research paper/preprint; and the recent FHIR work is a 2025 research preprint. ([Proceedings](#))

Below is the **final 30-reference set I recommend for MediGenie**.

Final 30 IEEE References for MediGenie

A. Dataset references — 1–9

These should be cited in your **Dataset / Data Sources** section. Do not replace dataset provenance with a random ML paper that happens to use the dataset.

[1] Heart Disease

A. Janosi, W. Steinbrunn, M. Pfisterer, and R. Detrano, “Heart Disease,” UCI Machine Learning Repository, 1989.

[Official UCI dataset / citation record](#)

Cite in: Section 3 — Dataset and Data Sources → Heart Disease.

The UCI record identifies the Cleveland/Heart Disease dataset and its original contributors. ([UCI Machine Learning Repository](#))

[2] Diabetes 130-US Hospitals

B. Strack et al., “Impact of HbA1c Measurement on Hospital Readmission Rates: Analysis of 70,000 Clinical Database Patient Records,” *BioMed Research International*, 2014.

Cite in: Section 3 — Dataset and Data Sources → Diabetes.

Important: use this citation to establish the **Diabetes 130-US Hospitals/readmission** dataset. In your paper, call the MediGenie task **Diabetes Patient Hospital Readmission Risk**, not diabetes diagnosis.

[3] Chronic Kidney Disease

L. Rubini, P. Soundarapandian, and P. Eswaran, “Chronic Kidney Disease,” UCI Machine Learning Repository, 2015.

[Official UCI dataset / DOI](#)

Cite in: Section 3 → Kidney Disease.

The UCI record gives the dataset citation and DOI 10.24432/C5G020. ([UCI Machine Learning Repository](#))

[4] Indian Liver Patient Dataset

B. Ramana and N. Venkateswarlu, “ILPD (Indian Liver Patient Dataset),” UCI Machine Learning Repository, 2022.

[Official UCI dataset / DOI](#)

Cite in: Section 3 → Liver Disease.

UCI currently gives this dataset citation and DOI. ([UCI Machine Learning Repository](#))

[5] Breast Cancer Wisconsin Diagnostic

W. Wolberg, O. Mangasarian, N. Street, and W. Street, “Breast Cancer Wisconsin (Diagnostic),” UCI Machine Learning Repository, 1993.

[Official UCI dataset / DOI](#)

Cite in: Section 3 → Breast Cancer.

UCI identifies this as the Wisconsin Diagnostic dataset and provides the official citation. ([UCI Machine Learning Repository](#))

[6] Parkinson's

H. A. Little et al., “Parkinsons,” UCI Machine Learning Repository, 2007.

[Official UCI Parkinsons dataset](#)

Cite in: Section 3 → Parkinson's Disease.

This is the dataset you actually use. **Do not cite Parkinson's Telemonitoring as your MediGenie dataset.**

[7] Hepatitis

R. S. Gong, “Hepatitis,” UCI Machine Learning Repository, 1988.

[Official UCI Hepatitis dataset](#)

Cite in: Section 3 → Hepatitis.

[8] Heart Failure Clinical Records

D. Chicco and G. Jurman, “Machine learning can predict survival of patients with heart failure from serum creatinine and ejection fraction alone,” *BMC Medical Informatics and Decision Making*, vol. 20, no. 16, 2020.

[DOI / official paper](#)

Cite in: Section 3 → Heart Failure and Section 6 → Comparison with Previous Research.

This is one of your **strongest direct comparison papers**, because it uses the Heart Failure Clinical Records dataset. ([PubMed Central \(PMC\)](#))

[9] Stroke Prediction Dataset

fedesoriano, “Stroke Prediction Dataset,” Kaggle, 2021.

[Official Kaggle dataset](#)

Cite in: Section 3 → Stroke.

The Kaggle record identifies the dataset, author and target variable. ([Kaggle](#))

Important: don't pretend this is a peer-reviewed paper. Cite it as a **dataset**.

B. Machine-learning methodology — 10–15

These support the actual algorithms and evaluation methodology in MediGenie.

[10] Random Forest

L. Breiman, “Random forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.

[DOI / Springer](#)

Cite in: Section 4 — Machine Learning Methodology → Random Forest.

The original paper is the correct foundational citation. ([DOI.org](#))

[11] Support Vector Machine

C. Cortes and V. Vapnik, “Support-vector networks,” *Machine Learning*, vol. 20, pp. 273–297, 1995.

[DOI / Springer](#)

Cite in: Section 4 → SVM.

([DOI.org](#))

[12] XGBoost

T. Chen and C. Guestrin, “XGBoost: A scalable tree boosting system,” in *Proc. 22nd ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining*, 2016, pp. 785–794.

[DOI / ACM](#)

Cite in: Section 4 → XGBoost.

([DOI.org](#))

[13] Logistic Regression

D. W. Hosmer, S. Lemeshow, and R. X. Sturdivant, *Applied Logistic Regression*, 3rd ed. Wiley, 2013.

Cite in: Section 4 → Logistic Regression.

[14] SMOTE / Class Imbalance

N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, “SMOTE: Synthetic minority over-sampling technique,” *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002.

Cite in: Section 4 → Class imbalance / preprocessing and Section 6 → limitations.

Use this when discussing why accuracy alone can be misleading on datasets such as Stroke and Hepatitis.

[15] Matthews Correlation Coefficient

D. Chicco, N. Tötsch, and G. Jurman, “The Matthews correlation coefficient (MCC) is more reliable than F1 score and accuracy in binary classification evaluation,” *BioData Mining*, vol. 14, 2021.

Cite in: Section 4 → Evaluation Metrics.

Use it to justify including MCC alongside accuracy, precision, recall, F1 and ROC-AUC.

C. Explainable AI — 16–18

[16] SHAP

S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” in *Advances in Neural Information Processing Systems*, 2017.

[Official paper / arXiv record](#)

Cite in: Section 4 → Explainability / SHAP.

The paper introduces SHAP as a unified framework for feature attribution. ([arXiv](#))

[17] XAI Healthcare Systematic Review

J. Jung, H. Lee, H. Jung, and H. Kim, “Essential properties and explanation effectiveness of explainable artificial intelligence in healthcare: A systematic review,” *Heliyon*, vol. 9, no. 5, 2023.

[DOI / PubMed](#)

Cite in: Introduction + Related Work + Explainability.

This is particularly useful for explaining why explainability matters in healthcare rather than simply saying “we used SHAP.” ([PubMed](#))

[18] Recent XAI/CDSS review

N. A. Aziz et al., “Explainable AI in Healthcare: Systematic Review of Clinical Decision Support Systems,” 2024.

[Official preprint record](#)

Cite in: Related Work → Explainable Clinical Decision Support.

This is a **preprint**, so label it appropriately in your bibliography rather than implying it is a journal article. ([MedRxiv](#))

D. RAG / Medical LLMs — 19–24

These are especially important because RAG + LLM reasoning is one of the major parts of MediGenie.

[19] Original RAG

P. Lewis et al., “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Advances in Neural Information Processing Systems*, 2020.

[Official research record](#)

Cite in: Section 4 → RAG Architecture.

The original work explicitly combines parametric generation with retrieved non-parametric knowledge. ([arXiv](#))

[20] Large Language Models in Medicine

A. J. Thirunavukarasu et al., “Large language models in medicine,” *Nature Medicine*, vol. 29, pp. 1930–1940, 2023.

[DOI / Nature](#)

Cite in: Introduction + Related Work → Medical LLMs.

This is one of the best broad medical-LLM references for your paper. ([Nature](#))

[21] Medical LLM Survey

H. Zhou et al., “A survey of large language models in medicine: Progress, application, and challenge,” 2023.

[Official research record](#)

Cite in: Related Work → Medical LLM landscape and limitations.

([arXiv](#))

[22] GPT-4 Medical Evaluation

H. Nori et al., “Capabilities of GPT-4 on medical challenge problems,” 2023.

[Microsoft Research / paper record](#)

Cite in: Related Work → Medical reasoning / LLM limitations.

The study evaluates GPT-4 on USMLE-style material and medical benchmarks and explicitly discusses calibration and safety. ([Microsoft](#))

[23] Med-PaLM 2

K. Singhal et al., “Towards expert-level medical question answering with large language models,” 2023.

[Official research record](#)

Cite in: Related Work → Medical question answering and clinical reasoning.

The paper reports Med-PaLM 2's medical QA performance and emphasizes the need for further real-world validation. ([arXiv](#))

[24] LLM Healthcare Systematic Review

L. Wang et al., “Applications and concerns of ChatGPT and other conversational large language models in health care: Systematic review,” *Journal of Medical Internet Research*, vol. 26, e22769, 2024.

[DOI / PubMed](#)

Cite in: Related Work + Discussion → Risks/limitations of medical LLMs.

([PubMed](#))

E. Multi-agent medical AI — 25–27

This is where your MediGenie architecture becomes much more interesting.

[25] MedAgents

X. Tang et al., “MedAgents: Large language models as collaborators for zero-shot medical reasoning,” 2023.

[Official research record](#)

Cite in: Related Work → Multi-agent medical reasoning.

MedAgents uses multiple LLM-based agents that independently analyze cases and collaborate toward a consensus. ([arXiv](#))

[26] MDAgents

Y. Kim et al., “MDAgents: An adaptive collaboration of LLMs for medical decision-making,” in *Advances in Neural Information Processing Systems* 37, 2024.

[NeurIPS DOI](#)

Cite in: Related Work → Multi-agent medical decision-making + Discussion.

This is one of the **most important recent references for MediGenie**. MDAgents evaluates adaptive collaboration structures for medical tasks and reports improvements across medical benchmarks. ([Proceedings](#))

[27] Medical LLM/Agent evaluation

S. Delourme et al., “Measured performance and healthcare professional perception of large language models used as clinical decision support systems: A scoping review,” 2024.

[DOI / IOS Press](#)

Cite in: Related Work → Clinical decision support + Discussion → Need for human oversight.

The review specifically addresses LLMs as clinical decision-support systems and healthcare-professional perception. ([Sage Journals](#))

F. FHIR / interoperability — 28–29

[28] FHIR foundation

K. Bender and K. Sartipi, “HL7 FHIR: An Agile and RESTful approach to healthcare information exchange,” in *Proc. IEEE International Conference on Healthcare Informatics*, 2013.

Cite in: System Architecture → Interoperability/FHIR.

Use this as the foundational FHIR citation.

[29] LLM + FHIR

A. Riquelme, P. Costa, and C. Martinez, “Large language models for automating clinical data standardization: HL7 FHIR use case,” 2025.

[Official research record](#)

Cite in: Related Work → FHIR + LLM interoperability and your FHIR implementation discussion.

The study specifically investigates LLM-based conversion of clinical data into HL7 FHIR resources and reports both high mapping performance and failure modes such as hallucinated attributes. ([arXiv](#))

Important: this is a **2025 preprint**, so label it as such.

G. Final recent clinical-AI reference — 30

[30] Biomedical/Health Informatics LLM review

H. Yu et al., “Large language models in biomedical and health informatics: A review with bibliometric analysis,” 2024.

[Official research record](#)

Cite in: Introduction + Related Work → Biomedical AI/LLM landscape.

This review analyzed 1,698 research articles and discusses clinical decision support, medical document analysis, patient interaction, and validation challenges. ([arXiv](#))

Where to put the references in your MediGenie paper

Here is the **exact citation map** I recommend.

Paper section	References
Abstract	Usually none
1. Introduction	[20], [24], [26], [30]
2. Related Work – ML disease prediction	[8], [10]–[15]

Paper section	References
2. Related Work – Medical LLMs	[20]–[24]
2. Related Work – Multi-agent AI	[25]–[27]
2. Related Work – RAG	[19], [20], [24]
2. Related Work – Explainable AI	[16]–[18]
2. Related Work – FHIR	[28], [29]
3. Dataset and Data Sources	[1]–[9]
4. ML Methodology	[10]–[15]
4. Explainability	[16]–[18]
4. RAG architecture	[19]
4. Multi-agent architecture	[25], [26]
4. FHIR interoperability	[28], [29]
5. Experimental Results	Your own results — cite methodology refs where appropriate, but do not cite another paper as evidence for your own numbers
6. Comparison with Previous Work	[8], [10]–[12], plus disease-specific literature
6. Discussion	[17], [20], [24], [26], [27], [29], [30]
7. Limitations	[18], [20], [24], [27], [29]
8. Conclusion	Usually no new references

One important correction to your comparison strategy

Your paper should **not** say:

"MediGenie is better than all previous research."

That's not supportable from your current experiments.

Instead, use:

"MediGenie provides a unified multi-disease clinical AI architecture that integrates disease-specific machine-learning prediction, explainability, multi-agent reasoning, retrieval-augmented knowledge access, and FHIR-based interoperability."

Then compare **model-by-model** where the dataset and task definition are genuinely comparable.

For example, your Heart Failure model can be compared with Chicco & Jurman's study because the UCI dataset is explicitly associated with that paper. ([UCI Machine Learning Repository](#))

For Stroke, be more careful: your source is the Kaggle Stroke Prediction Dataset, so don't claim direct reproduction of another paper unless that paper demonstrably used the **same dataset/version and target definition**. The Kaggle record identifies the exact dataset you use. ([Kaggle](#))

The 5 references I consider essential for your novelty claim

If your examiner asks, "**What is new about MediGenie?**", these are the five I would immediately discuss:

- [19] RAG — retrieval-grounded knowledge
- [25] MedAgents — collaborative medical agents
- [26] MDAgents — adaptive multi-agent medical decision-making
- [29] LLM + FHIR interoperability
- [17] Explainable AI in healthcare

Together they support the architecture around your nine ML predictors rather than pretending that SVM/RF/XGBoost themselves are novel. MDAgents is particularly strong because it is a recent NeurIPS paper directly addressing collaborative LLM-based medical decision-making. ([Proceedings](#))

One final rule

Do not add a reference to your bibliography unless you actually cite it somewhere in the paper. Conversely, every [1]–[30] reference above should appear at least once in the manuscript.

Also distinguish **peer-reviewed papers**, **preprints**, and **datasets** in your bibliography. In particular, [18], [21]–[23], [25], [29], and [30] include research/preprint records rather than all being conventional journal publications. This keeps your "30 real references" list academically honest rather than making a preprint look like a peer-reviewed journal article.

REFERENCES

- [1] A. Janosi, W. Steinbrunn, M. Pfisterer, and R. Detrano,
“Heart Disease,” UCI Machine Learning Repository, 1989.
[Online]. Available: <https://archive.ics.uci.edu/dataset/45/heart+disease>
- [2] B. Strack et al.,
“Impact of HbA1c Measurement on Hospital Readmission Rates:
Analysis of 70,000 Clinical Database Patient Records,”
BioMed Research International, vol. 2014, Art. no. 781670, 2014.
- [3] L. Rubini, P. Soundarapandian, and P. Eswaran,
“Chronic Kidney Disease,” UCI Machine Learning Repository, 2015.
[Online]. Available: <https://archive.ics.uci.edu/dataset/336/chronic+kidney+disease>
- [4] B. Ramana and N. Venkateswarlu,
“ILPD (Indian Liver Patient Dataset),”
UCI Machine Learning Repository, 2012.
[Online]. Available:
<https://archive.ics.uci.edu/dataset/225/ilpd+indian+liver+patient+dataset>
- [5] W. H. Wolberg, O. L. Mangasarian, N. Street, and W. H. Street,
“Breast Cancer Wisconsin (Diagnostic),”
UCI Machine Learning Repository, 1995.
[Online]. Available:
<https://archive.ics.uci.edu/dataset/17/breast+cancer+wisconsin+diagnostic>
- [6] H. A. Little et al.,
“Parkinsons,” UCI Machine Learning Repository, 2007.
[Online]. Available:

<https://archive.ics.uci.edu/dataset/174/parkinsons>

[7] R. S. Gong,

“Hepatitis,” UCI Machine Learning Repository, 1988.

[Online]. Available:

<https://archive.ics.uci.edu/dataset/46/hepatitis>

[8] D. Chicco and G. Jurman,

“Machine learning can predict survival of patients with heart failure from serum creatinine and ejection fraction alone,”

BMC Medical Informatics and Decision Making, vol. 20,

Art. no. 16, 2020, doi: 10.1186/s12911-020-1023-5.

[9] fedesoriano,

“Stroke Prediction Dataset,” Kaggle, 2021.

[Online]. Available:

<https://www.kaggle.com/datasets/fedesoriano/stroke-prediction-dataset>

[10] L. Breiman,

“Random forests,” Machine Learning, vol. 45, no. 1,

pp. 5–32, 2001, doi: 10.1023/A:1010933404324.

[11] C. Cortes and V. Vapnik,

“Support-vector networks,” Machine Learning, vol. 20,

pp. 273–297, 1995, doi: 10.1007/BF00994018.

[12] T. Chen and C. Guestrin,

“XGBoost: A scalable tree boosting system,”

in Proc. 22nd ACM SIGKDD Int. Conf. Knowledge Discovery and Data Mining,

2016, pp. 785–794, doi: 10.1145/2939672.2939785.

[13] D. W. Hosmer, S. Lemeshow, and R. X. Sturdivant,
Applied Logistic Regression, 3rd ed.
Hoboken, NJ, USA: Wiley, 2013.

[14] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer,
“SMOTE: Synthetic minority over-sampling technique,”
Journal of Artificial Intelligence Research, vol. 16,
pp. 321–357, 2002.

[15] D. Chicco, N. Tötsch, and G. Jurman,
“The Matthews correlation coefficient (MCC) is more reliable
than F1 score and accuracy in binary classification evaluation,”
BioData Mining, vol. 14, Art. no. 13, 2021.

[16] S. M. Lundberg and S.-I. Lee,
“A unified approach to interpreting model predictions,”
in Advances in Neural Information Processing Systems,
vol. 30, 2017.

[17] J. Jung, H. Lee, H. Jung, and H. Kim,
“Essential properties and explanation effectiveness of explainable
artificial intelligence in healthcare: A systematic review,”
Heliyon, vol. 9, no. 5, Art. no. e16110, 2023.

[18] N. A. Aziz et al.,
“Explainable AI in healthcare: Systematic review of clinical
decision support systems,” medRxiv, 2024.

[Online]. Available:

<https://doi.org/10.1101/2024.08.10.24311735>

[19] P. Lewis et al.,

“Retrieval-augmented generation for knowledge-intensive NLP tasks,”

in *Advances in Neural Information Processing Systems*,

vol. 33, 2020.

[20] A. J. Thirunavukarasu et al.,

“Large language models in medicine,”

Nature Medicine, vol. 29, pp. 1930–1940, 2023,

doi: 10.1038/s41591-023-02448-8.

[21] H. Zhou et al.,

“A survey of large language models in medicine:

Progress, application, and challenge,” 2023,

arXiv:2311.05112.

[22] H. Nori et al.,

“Capabilities of GPT-4 on medical challenge problems,”

Microsoft Research, 2023.

[Online]. Available:

<https://www.microsoft.com/en-us/research/publication/capabilities-of-gpt-4-on-medical-challenge-problems/>

[23] K. Singhal et al.,

“Towards expert-level medical question answering

with large language models,” 2023,

arXiv:2305.09617.

[24] L. Wang et al.,
“Applications and concerns of ChatGPT and other conversational
large language models in health care: Systematic review,”
Journal of Medical Internet Research, vol. 26, 2024.

[25] X. Tang, A. Zou, Z. Zhang, Z. Li, Y. Zhao, X. Zhang,
A. Cohan, and M. Gerstein,
“MedAgents: Large language models as collaborators
for zero-shot medical reasoning,”
in Findings of the Association for Computational Linguistics:
ACL 2024, Bangkok, Thailand, 2024, pp. 599–621,
doi: 10.18653/v1/2024.findings-acl.33.

[26] Y. Kim et al.,
“MDAgents: An adaptive collaboration of LLMs
for medical decision-making,”
in Advances in Neural Information Processing Systems 37,
2024, pp. 79410–79452,
doi: 10.52202/079017-2522.

[27] S. Delourme et al.,
“Measured performance and healthcare professional perception
of large language models used as clinical decision support systems:
A scoping review,”
Studies in Health Technology and Informatics, 2024,
doi: 10.3233/SHTI240543.

- [28] K. Bender and K. Sartipi,
“HL7 FHIR: An agile and RESTful approach to healthcare
information exchange,”
in Proc. IEEE International Conference on Healthcare Informatics,
2013.
- [29] A. Riquelme, P. Costa, and C. Martinez,
“Large language models for automating clinical data standardization:
HL7 FHIR use case,” 2025, arXiv:2507.03067.
- [30] H. Yu et al.,
“Large language models in biomedical and health informatics:
A review with bibliometric analysis,” 2024,
arXiv:2403.16303.